
Low-dimensional Interpretable Kernels with Conic Discriminant Functions for Classification

Gürhan Ceylan*

Eskişehir Teknik Üniversitesi
Eskişehir, Turkey
gurhanceylan@eskisehir.edu.tr

Ş. İlker Birbil

Econometric Institute
Erasmus University Rotterdam
3000 DR Rotterdam, The Netherlands
birbil@ese.eur.nl

Abstract

Kernels are often developed and used as implicit mapping functions that show impressive predictive power due to their high-dimensional feature space representations. In this study, we gradually construct a series of simple feature maps that lead to a collection of interpretable low-dimensional kernels. At each step, we keep the original features and make sure that the increase in the dimension of input data is extremely low, so that the resulting discriminant functions remain interpretable and amenable to fast training. Despite our persistence on interpretability, we obtain high accuracy results even without in-depth hyperparameter tuning. Comparison of our results against several well-known kernels on benchmark datasets show that the proposed kernels are competitive in terms of prediction accuracy, while the training times are significantly lower than those obtained with state-of-the-art kernel implementations.

1 Introduction

Whenever the relationship between input and output of a dataset is not linear, kernels are indispensable tools to model this nonlinearity by expanding the feature space. This expansion is sometimes carried with an explicit function called the *feature map*. Capturing the nonlinear relationship has been observed countless times to have a profound effect on the estimation accuracy. However, using kernels has two important caveats from a practitioner’s point of view: First, the dimension of the problem can become quite large after expanding the feature space. In fact for some of the kernels, even the explicit feature map is not known. Thus, the interpretation of the trained models becomes a daunting task. Second, the increase in the problem dimension causes an increase in the training times. When explicit feature map is known, then fast linear algorithms can be used [4, 16]. However, even in this case, expanding the feature space to large dimensions can still slow down the overall training process. In this study, we propose several simple feature maps that lead to a collection of interpretable kernels with varying degrees of freedom. We make sure that the increase in the dimension of input data with each proposed feature map is extremely low, so that the resulting models can be trained quickly, and the obtained results can easily be interpreted.

Approximating kernels via feature maps is a common approach to achieve fast training [16, 12, 11, 18, 15, 9, 1, 17]. Generally speaking, the success of an approximation increases as the dimension of the feature map increases. Therefore, a practitioner needs to conduct a series of experiments to find a feature map with a desirable performance in terms of accuracy and training time. When dimension increases, the resulting model becomes more difficult to interpret. For instance, the approximation method proposed by Pham and Pagh [15] achieves the same classification accuracy of a second order polynomial kernel by an almost 10-fold increase in the dimension. Our work has also ties with

*Corresponding author.

piecewise linear mapping functions introduced by Huang et al. [10]. However, like approximation methods, their results depend on the choice of the dimension. The authors report numerical results where the increase in the dimension with the proposed feature maps goes up to 20-fold.

In this study, we center our contributions around two schemes to construct low-dimensional feature maps. The first scheme increments the dimension of the input data only by one, whereas the second scheme doubles dimension of the feature map. These dimensions are much lower than the existing feature maps in the literature. Our feature maps may require an anchor point as a hyperparameter. We argue that this hyperparameter can either be set by the domain experts or by the training algorithms automatically. Combined with the low dimensionality, the results obtained with the proposed maps can be easily interpreted by the practitioners. We also discuss several conditions under which perfect separability is guaranteed for binary classification datasets. This observation can also be used to propose other methods for selecting anchor points. To elaborate on these points, we reserve a section in the supplementary document about how to extend the proposed schemes to intermediate dimensions as well as to multi-class classification.

In a different line of work, Gasimov and Ozturk [8] and Cevikalp and Triggs [3] introduce classifiers based on polyhedral conic functions that are similar to our feature maps. Both papers have used the polyhedral conic functions within optimization problems but have not established the relations of these functions to feature maps and their associated kernels. In a recent follow-up study, Cevikalp and Saglamilar [2] have also mentioned the ellipsoidal conic functions. Here, we consider a general class that includes both the polyhedral and the ellipsoidal conic functions. We also explicitly present the resulting feature maps along with the corresponding kernels. This shows that the related classifiers in the literature are just kernel methods. Our discussion through explicit feature maps has far reaching consequences, as associated kernels can be used in various learning methods.

In the light of this review, we make the following contributions to the literature:

- We propose several low-dimensional feature maps, which simply concatenates the original input features with distance-based features. The proposed feature maps are easy to interpret and their training times are in par with linear kernels.
- We compare the new kernels against the commonly used kernels on a set of binary classification datasets. Our numerical results demonstrate that the proposed kernels obtain high accuracy, and on several datasets, even outperform all the other kernels. We also note that our results with the proposed kernels are obtained without any parameter tuning.
- We demonstrate with our implementation² that it is truly simple to incorporate the proposed kernels within the existing software packages used for kernel-based methods.

2 Low-dimensional feature maps and kernels

Consider a practitioner, who trains a classification model on a dataset to obtain the weights associated with the features of the d -dimensional input vector \mathbf{x} . When the relationship is linear, we obtain the most interpretable model with the discriminant function³

$$f(\mathbf{x}) = \mathbf{w}_{1:d}^T \mathbf{x} = w_1 x_1 + w_2 x_2 + \dots + w_d x_d,$$

where $\mathbf{w}_{1:d}$ stands for the d -dimensional weight vector. Our first model simply asks the practitioner to set an anchor point \mathbf{a} and measure its distance to the input. That is

$$f_{p,1}(\mathbf{x} | \mathbf{a}) = \mathbf{w}_{1:d+1}^T \phi_{p,1}(\mathbf{x} | \mathbf{a}) = \mathbf{w}_{1:d}^T \mathbf{x} + w_{d+1} \|\mathbf{x} - \mathbf{a}\|_p^p, \quad (1)$$

where $\phi_{p,1}(\mathbf{x} | \mathbf{a}) : \mathbb{R}^d \mapsto \mathbb{R}^{d+1}$ with $p > 0$ is the proposed feature map that adds *only one* dimension, and w_{d+1} is the weight corresponding to the new feature measuring the distance between the anchor and the input⁴. Explicitly, this feature map is given by

$$\phi_{p,1}(\mathbf{x} | \mathbf{a}) = (x_1, x_2, \dots, x_d, \|\mathbf{x} - \mathbf{a}\|_p^p)^T = (\mathbf{x}^T, \|\mathbf{x} - \mathbf{a}\|_p^p)^T.$$

Throughout our discussion, we mainly use either $p = 1$ or $p = 2$, which are the two most common choices in a wide-range of learning algorithms. The discriminant function obtained with this feature

²(GitHub page) – Please see the supplementary file for the implementation.

³We omit the bias term added to the discriminant function to simplify our exposition.

⁴In case $p = \infty$, we abuse the notation slightly and work with ℓ_∞ -norm.

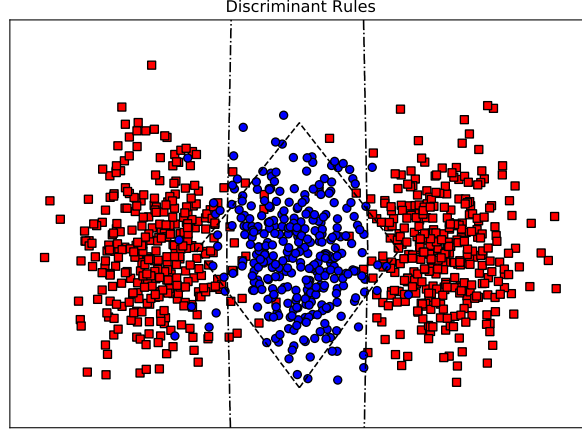


Figure 1: Conic discriminant functions obtained with a single anchor point for a binary classification dataset. Here, the anchor point is taken as the average of all samples.

map is quite stringent in the sense that only one more weight (w_{d+1}) is added with the anchor point. Thus, one immediate extension could be using different weights for each feature. In fact, this extension takes us to our second feature map $\phi_{p,d}(\mathbf{x}) : \mathbb{R}^d \mapsto \mathbb{R}^{2d}$ and the corresponding discriminant function. For $p = 1$, we obtain

$$f_{1,d}(\mathbf{x} | \mathbf{a}) = \mathbf{w}_{1:2d}^\top \phi_{1,d}(\mathbf{x} | \mathbf{a}) = \mathbf{w}_{1:d}^\top \mathbf{x} + \sum_{\ell=1}^d w_{d+\ell} |x_\ell - a_\ell|.$$

If we further define a diagonal matrix \mathbf{W} with elements $w_{d+1}, w_{d+2} \dots, w_{2d}$, then the discriminant function for $p = 2$ with the second feature map becomes

$$f_{2,d}(\mathbf{x} | \mathbf{a}) = \mathbf{w}_{1:2d}^\top \phi_{2,d}(\mathbf{x} | \mathbf{a}) = \mathbf{w}_{1:d}^\top \mathbf{x} + (\mathbf{x} - \mathbf{a})^\top \mathbf{W} (\mathbf{x} - \mathbf{a}).$$

Contrasting this discriminant function with (1) shows that after doubling the dimension, we still measure a *distance* to the anchor point. When the elements of \mathbf{W} are restricted to be positive, then the last term indeed becomes a weighted norm. The feature map that doubles the dimension is given by

$$\phi_{p,d}(\mathbf{x} | \mathbf{a}) = (x_1, x_2, \dots, x_d, |x_1 - a_1|^p, \dots, |x_d - a_d|^p)^\top.$$

Figure 1 illustrates two discriminant rules obtained with feature maps $\phi_{1,1}$ (diamond) and $\phi_{1,d}$ (vertical lines) for a two-dimensional binary classification problem. The flexibility of using more dimensions with $\phi_{1,d}$ provides a clear separation for classification. Note that the regions defined by the rules are the lower-level sets of different cones pointed at the anchor point. Thus, we refer to the proposed functions as *conic discriminant functions*. These conic discriminant functions have ties with the line of work initiated by Gasimov and Ozturk [8] and then extended by Cevikalp and Triggs [3]. In these two works, the discriminant functions are coined as *polyhedral conic classifiers*. In Section S.3, we elaborate on the relationship between our current work and the conic classifiers.

Up to this point, we have not discussed how to select the anchor point \mathbf{a} . Consider a dataset consisting of the samples $\mathbf{x}_i \in \mathbb{R}^d$, $i \in \mathcal{I} = \{1, \dots, m\}$. If the user prefers an automatic selection for the anchor point, then one straightforward choice is the average of all samples. Depending on the application, the anchor point could as well be decided by the domain experts. Likewise, the experts may also propose a set of anchor points \mathcal{A} instead of just one point. Given the set of anchor point \mathcal{A} , we may select the closest point to sample \mathbf{x}_i as its anchor by evaluating

$$\mathbf{a}_i = \arg \min_{\mathbf{a} \in \mathcal{A}} \{\|\mathbf{x}_i - \mathbf{a}\|_p\}. \quad (2)$$

We can then map \mathbf{x}_i to a higher-dimensional space with $\phi_{p,1}(\mathbf{x}_i | \mathbf{a}_i)$ or $\phi_{p,d}(\mathbf{x}_i | \mathbf{a}_i)$. Figure 2 shows the increase in the dimension as well as the obtained discriminant rules for two binary classification problems. The top row is given for a single anchor point, which is taken as the average of all samples. The bottom row in the same figure shows the mappings when the set of anchor points

\mathcal{A} is provided by the user. Here, \mathcal{A} is constructed with the sample averages of the clusters of one class (blue circles). Then, the anchor point for each sample is selected by applying (2). The set \mathcal{A} can also be constructed when different samples are known to be associated with different clusters. These clusters may already shape during data collection for instance when multiple cohorts, spatial differences, temporal variations, and so on, are involved. Clearly, different clusters may as well be formed algorithmically beforehand by using unsupervised learning methods.

In the main text of our following discussion, we will use a single anchor point. We have reserved Section S.2⁵) to introduce different examples, where multiple anchor points are selected. We will also consider mostly binary classification problems, since kernels are frequently used within the well-known Support Vector Machine (SVM) algorithm. The following proposition formally shows the conditions, under which a linearly nonseparable binary classification problem can be mapped to a linearly separable feature space with a single anchor point using the proposed feature maps.

Proposition 1 *Suppose we have a linearly nonseparable dataset with samples $\mathbf{x}_i \in \mathbb{R}^d$ and the corresponding labels $y_i \in \{+1, -1\}$ for $i \in \mathcal{I}$. If we further define two index sets $\mathcal{I}^+ = \{i \in \mathcal{I} : y_i = +1\}$ and $\mathcal{I}^- = \{i \in \mathcal{I} : y_i = -1\}$, then using $\phi_{p,1}(\mathbf{x} | \mathbf{a})$ returns a linearly separable dataset, if the chosen anchor point satisfies*

$$\min_{i \in \mathcal{I}^+} \{\|\mathbf{x}_i - \mathbf{a}\|_p^p\} > \max_{i \in \mathcal{I}^-} \{\|\mathbf{x}_i - \mathbf{a}\|_p^p\} \text{ or } \max_{i \in \mathcal{I}^+} \{\|\mathbf{x}_i - \mathbf{a}\|_p^p\} < \min_{i \in \mathcal{I}^-} \{\|\mathbf{x}_i - \mathbf{a}\|_p^p\}. \quad (3)$$

Likewise, using $\phi_{p,d}(\mathbf{x} | \mathbf{a})$ returns a linearly separable dataset if the chosen anchor point for some dimension $\ell \in \{1, \dots, d\}$ satisfies

$$\min_{i \in \mathcal{I}^+} \{|x_{i\ell} - a_\ell|^p\} > \max_{i \in \mathcal{I}^-} \{|x_{i\ell} - a_\ell|^p\} \text{ or } \max_{i \in \mathcal{I}^+} \{|x_{i\ell} - a_\ell|^p\} < \min_{i \in \mathcal{I}^-} \{|x_{i\ell} - a_\ell|^p\}. \quad (4)$$

The proof of this proposition is given in Section S.1. Although it is a straightforward result, Proposition 1 provides a clear point of view for the role of the anchor points. This view actually allows us to conduct a systematic search for selecting the anchor point. We discuss in Section S.2 one such approach yielding complex discriminant rules particularly useful for multi-class classification. In any case, we should point out that our numerical study on benchmark datasets in the next section show that the proposed feature maps achieve competitive classification accuracies even without explicitly searching for the *best* anchor point. We use the sample averages as the anchor points. Even further, we scale the data so that the sample average becomes the origin. Hence, the anchor point simply disappears.

We are ready to present the kernels associated with the proposed feature maps. Given samples \mathbf{x} and \mathbf{z} , the proposed kernel functions are obtained by simply taking the inner product in higher-dimensional space

$$k_{p,1}(\mathbf{x}, \mathbf{z} | \mathbf{a}) = \phi_{p,1}(\mathbf{x} | \mathbf{a})^\top \phi_{p,1}(\mathbf{z} | \mathbf{a}) \text{ and } k_{p,d}(\mathbf{x}, \mathbf{z} | \mathbf{a}) = \phi_{p,d}(\mathbf{x} | \mathbf{a})^\top \phi_{p,d}(\mathbf{z} | \mathbf{a}).$$

Let us now contrast our kernel functions to the well-known radial basis function and the polynomial kernel function given by

$$k_\gamma^{\text{RBF}}(\mathbf{x}, \mathbf{z}) = e^{-\gamma \|\mathbf{x} - \mathbf{z}\|_2^2}, \gamma > 0 \text{ and } k_q^{\text{POL}}(\mathbf{x}, \mathbf{z}) = (\mathbf{x}^\top \mathbf{z} + 1)^q, q \in \mathbb{N},$$

respectively. There is no explicit feature map for RBF, since the input data is mapped to an infinite dimensional space. Although, the polynomial kernel function is associated with a finite dimensional space, the resulting dimension can be quite large. Even for 20-dimensional input data, the feature space has more than 200 dimensions for $q = 2$. In addition to concerns about interpretability, these kernels may also slow down the training process. Suppose that we train a SVM model on the samples $\mathbf{x}_i \in \mathbb{R}^d$ with the class labels $y_i \in \{-1, +1\}$ for $i \in \mathcal{I}$. The training requires to store an $m \times m$ matrix consisting of kernel function evaluations of all sample pairs. Then, a linear system is solved with this matrix. Thus, the computation time complexity is in the order of $\mathcal{O}(m^2)$ to $\mathcal{O}(m^3)$. After training, the discriminant rule of the SVM model becomes

$$f(\mathbf{x}) = \sum_{i=1}^m \alpha_i y_i k(\mathbf{x}_i, \mathbf{x})$$

where α_i values are the weights obtained with training. If the explicit feature map ϕ_\bullet is known, then we can also obtain the weights with $\mathbf{w} = \sum_{i=1}^m \alpha_i y_i \phi_\bullet(\mathbf{x}_i)$.

⁵All cross references starting with “S” refer to the supplementary document.

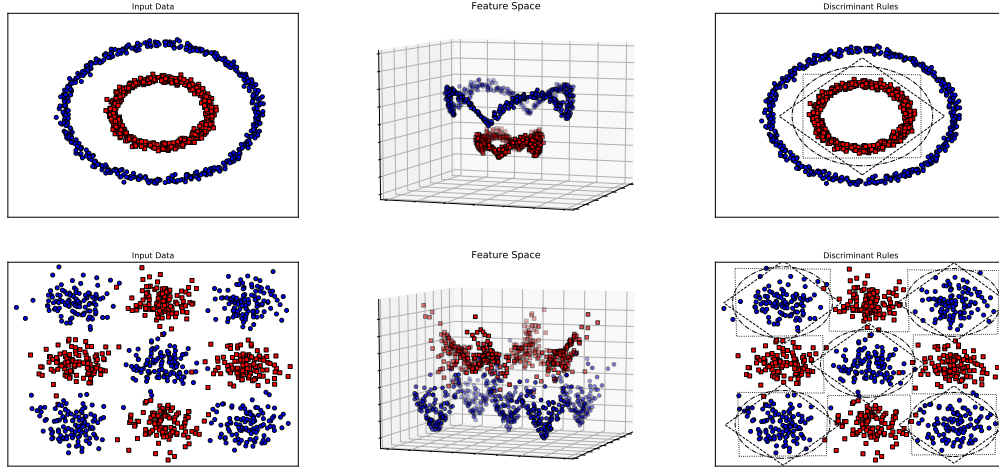


Figure 2: Visualization of selecting single (top row) or multiple (bottom row) anchor points. The feature spaces are obtained with $p = 1$. The discriminant rules are given for $p = 1$ (diamond), $p = 2$ (ellipsoid) and $p = \infty$ (rectangular). The anchor point in the top row is simply the average of the whole dataset. The set of anchor points in the bottom row are the sample averages within the clusters of blue circles.

Table 1: Average accuracies obtained with different methods on small datasets

Datasets	LIN	$\phi_{1,1}$	$\phi_{2,1}$	$\phi_{1,d}$	$\phi_{2,d}$	POL	RBF
Australian	86.81	86.81	86.67	86.67	86.09	84.20	85.51
Fourclass	76.21	78.07	79.35	77.14	79.70	79.36	100.0
Ionosphere	87.72	91.44	92.57	91.44	91.73	91.44	93.72
Heart	84.07	84.44	84.07	84.81	83.33	82.96	84.44
Pima	77.34	76.82	76.68	76.69	76.95	73.43	76.82
W.Prognostic	81.37	79.84	79.34	80.37	77.32	76.32	77.74
Bupa	69.77	69.5	69.18	73.62	73.01	61.85	72.98
Fertility	87.00	87.00	87.00	87.00	88.00	86.00	86.00
W.Diagnostic	97.07	96.93	97.22	96.78	96.78	95.61	96.93

3 Numerical experiments

In this section, we first compare the performances of the proposed kernels against linear (LIN), radial basis function (RBF) and polynomial (POL) kernels on a set of binary classification datasets compiled from the literature [7, 4]. A summary of the datasets is given in Section S.4. This benchmarking study is conducted in terms of prediction accuracy and training time. We also demonstrate on a particular dataset that our feature maps constitute a clear choice when compared against two well-known approximation methods in terms of interpretability.

We have implemented all our kernels in Python⁶ using the scikit-learn [14] package. Since our feature maps are explicit, we have used the fast linear solver provided with this package. Thus, we also give the times spent for mappings to give a fair comparison. For small datasets, we have used 10-fold stratified cross validation and reported the averages. For the remaining ones, we have mostly used the test samples provided with the dataset. If a test set is not available, then we have applied a standard train (70%)-test (30%) split. In all our tests, the anchor point is used as the sample mean. Thus, scaling the dataset has allowed us to take $\mathbf{a} = \mathbf{0}$. The other hyperparameters are selected from the following sets: $\gamma \in \{10^i, i = -5, \dots, 4\}$, $q \in \{2, 3, 4\}$, $C \in \{10^i, i = -5, \dots, 4\}$. Here C is the SVM regularization parameter. The best performing parameters are chosen by applying grid search with stratified two-fold cross validation.

⁶(GitHub page) – Please see the supplementary file to reproduce all our results here.

Table 2: Accuracies obtained with different methods on large datasets

Datasets	LIN	$\phi_{1,1}$	$\phi_{2,1}$	$\phi_{1,d}$	$\phi_{2,d}$	POL	RBF
Splice	85.29	86.02	86.02	91.95	89.61	85.61	89.93
Wilt	70.60	83.60	81.20	85.60	84.00	84.00	81.80
Guide1	95.62	96.25	96.10	96.48	96.15	96.70	96.62
Spambase	92.76	92.25	92.90	94.93	93.05	91.89	93.70
Phoneme	75.46	73.61	74.29	76.82	76.57	78.36	87.55
Magic	79.43	80.98	80.30	85.61	84.37	84.40	87.71
Adult	84.93	84.93	84.94	84.93	84.92	84.39	85.06

Table 3: Training times in seconds

Datasets/Kernels	LIN	$\phi_{1,1}$	$\phi_{2,1}$	$\phi_{1,d}$	$\phi_{2,d}$	POL	RBF
Splice	<1	<1	<1	<1	<1	<1	<1
Wilt	<1	<1	<1	<1	<1	<1	<1
Guide1	<1	<1	<1	<1	<1	<1	<1
Spambase	<1	<1	<1	<1	1.89	<1	<1
Phoneme	<1	<1	<1	<1	<1	169.88	<1
Magic	<1	<1	<1	<1	<1	425.39	63.94
Adult	<1	1.20	1.21	1.39	7.43	89.20	151.36

Table 1 shows our results on small datasets, where the accuracy of the best performing method for the corresponding dataset is written in boldface. Except on Pima and W.Prognostic, the proposed kernels achieve better average accuracy values than the linear kernel. For the same two datasets, it is important to note that POL and RBF also perform worse than the linear kernel. Overall, POL does not achieve the best accuracy value for any one of the datasets. Though RBF is the clear winner in two datasets, it is outperformed by one of our kernels in all other problems. In fact, the proposed kernels achieve the best predictions on four of the datasets. Since these are small datasets, the training times of all kernels are negligible, and hence, we do not report those figures.

We report the accuracy values and the training times for large datasets in Table (2) and Table (3), respectively. Recall that the training times of the proposed kernels also include the time spent for explicit mappings. For the large datasets, we observe that the performance of the linear kernel degrades. Even if we use only one additional feature with $\phi_{1,1}$ and $\phi_{2,1}$, we obtain better accuracy values than LIN for all problems but one. For Wilt dataset, this increase goes up to 13% with almost no increase in training time. Both POL and RBF also return good accuracy values than LIN for large datasets. However, Table 3 shows that this improvement comes at a great cost in training times. For instance, the accuracy improvement achieved by POL and RBF on Magic dataset are around 5% and 8%, respectively. The corresponding training times are approximately 60 times and 425 times worse than the linear kernel. All our kernels, on the other hand, achieve accuracy values on par with POL and RBF in less than a second.




Figure 3: Accuracy and training time comparison against approximation methods on Magic dataset.

In the literature, kernel approximation methods are also used to obtain feature maps with varying dimensions. With these approximation methods, the accuracy and the computation times oftentimes increase as the dimension increases. Figure 3 shows an example comparison of our kernels against two well-known approximation approaches, Fourier transform method [16] and Nyström method [19]. In these experiments, we have used the setup for Magic dataset as above with $C = 1$ and $\gamma = 0.2$. A practitioner may consider to use the approximation methods to obtain low (transformed) dimensions for possible interpretability. However, as the left plot in Figure 3 shows, the classification accuracies of both approximation methods are poor for low transformed feature dimensions. Although an increase in the dimension increases the accuracy, the training time also goes up. Our first kernel $\phi_{2,1}$, which adds only one more feature, performs slightly better than the linear kernel. On the other hand, when use our second kernel $\phi_{1,d}$ and double the number of features dimensions, we obtain a classification accuracy value that is close to RBF. The right plot in Figure 3 shows that the training times of our kernels are almost the same as the linear kernel.

4 Conclusion

We propose several low-dimensional kernels that are derived from two groups of explicit feature maps. The first group expands the feature space of the samples by adding only one more feature and the second group doubles the number of features. Both groups append additional distance-based features to the original list of features. This limited expansion of the dimension with explicit feature maps brings two advantages for practitioners: interpretability and fast training. With our experimental results on several datasets, we have shown that these advantages do not come at the cost of accuracy. On the contrary, for several problems, we have obtained better accuracy values than those of the well-known kernels. Moreover, the training times of our feature maps have been only a small fraction of the times spent with those kernels.

The proposed distance-based feature maps are defined with respect to a single anchor point. In practice, this anchor point can be selected by the domain experts from a set of candidate points. We also argue that preprocessing of the samples with clustering methods may also be used to determine the anchor point. In this work, we have removed the anchor point from our computational study by scaling the datasets and taking their sample means as the anchor point. Our numerical results have shown that even without fine-tuning the selection of the anchor point, the proposed low-dimensional kernels perform remarkably well. Nonetheless, we believe that selection of one or more anchor points can play an important role in prediction accuracy. As we present in Section S.2, our initial experiments show that the accuracy values can be improved with multiple anchor points (leading to intermediate dimensions) at almost no cost of additional training time. Naturally, the resulting feature maps remain interpretable. The same line of reasoning can also lead to a systematic search of a set of anchor points for multi-class classification. These observations shape our future research agenda.

References

- [1] H. Avron, V. Sindhwani, J. Yang, and M. W. Mahoney. Quasi-monte carlo feature maps for shift-invariant kernels. *The Journal of Machine Learning Research*, 17(1):4096–4133, 2016.
- [2] H. Cevikalp and H. Saglamilar. Polyhedral conic classifiers for computer vision applications and open set recognition. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 2019.
- [3] H. Cevikalp and B. Triggs. Polyhedral conic classifiers for visual object detection and classification. In *The IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, July 2017.
- [4] C.-C. Chang and C.-J. Lin. Libsvm: A library for support vector machines. *ACM Transactions on Intelligent Systems and Technology (TIST)*, 2(3):1–27, 2011.
- [5] E. Cimen and G. Ozturk. O-pcf algorithm for one-class classification. *Optimization Methods and Software*, pages 1–15, 2019.
- [6] E. Cimen, G. Ozturk, and O. N. Gerek. Incremental conic functions algorithm for large scale classification problems. *Digital Signal Processing*, 77:187–194, 2018.
- [7] D. Dua and C. Graff. UCI machine learning repository, 2017. URL <http://archive.ics.uci.edu/ml>.
- [8] R. N. Gasimov and G. Ozturk. Separation via polyhedral conic functions. *Optimization Methods and Software*, 21(4):527–540, 2006.
- [9] R. Hamid, Y. Xiao, A. Gittens, and D. DeCoste. Compact random feature maps. In *International Conference on Machine Learning*, pages 19–27, 2014.
- [10] X. Huang, S. Mehrkanoon, and J. A. Suykens. Support vector machines with piecewise linear feature mapping. *Neurocomputing*, 117:118 – 127, 2013. ISSN 0925-2312.
- [11] P. Kar and H. Karnick. Random feature maps for dot product kernels. In *Artificial Intelligence and Statistics*, pages 583–591, 2012.
- [12] S. Maji and A. C. Berg. Max-margin additive classifiers for detection. In *2009 IEEE 12th International Conference on Computer Vision*, pages 40–47. IEEE, 2009.
- [13] G. Ozturk and M. T. Ciftci. Clustering based polyhedral conic functions algorithm in classification. *Journal of Industrial & Management Optimization*, 11(3):921, 2015.
- [14] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cournapeau, M. Brucher, M. Perrot, and E. Duchesnay. Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12:2825–2830, 2011.
- [15] N. Pham and R. Pagh. Fast and scalable polynomial kernels via explicit feature maps. In *Proceedings of the 19th ACM SIGKDD International Conference on Knowledge Discovery and Data mining*, pages 239–247, 2013.
- [16] A. Rahimi and B. Recht. Random features for large-scale kernel machines. In *Advances in Neural Information Processing Systems*, pages 1177–1184, 2008.
- [17] S. Si, C.-J. Hsieh, and I. S. Dhillon. Memory efficient kernel approximation. *The Journal of Machine Learning Research*, 18(1):682–713, 2017.
- [18] S. Vempati, A. Vedaldi, A. Zisserman, and C. Jawahar. Generalized RBF feature maps for efficient detection. In *BMVC*, pages 1–11, 2010.
- [19] T. Yang, Y.-F. Li, M. Mahdavi, R. Jin, and Z.-H. Zhou. Nyström method vs random fourier features: A theoretical and empirical comparison. In *Advances in Neural Information Processing Systems*, pages 476–484, 2012.

Supplementary Material for

“Low-dimensional Interpretable Kernels with Conic Discriminant Functions for Classification”

Note that most of the cross references in this supplementary file refer to the original manuscript. Therefore, clicking on those references may not take you to the desired location. All the results and the figures both in the main document and in this supplementary document can be reproduced with the scripts that are provided in the accompanying compressed file.

S.1 Omitted proof

PROOF OF PROPOSITION 1: Suppose that the first part of (3) holds and define

$$\mu = \min_{i \in \mathcal{I}^+} \{\|\mathbf{x}_i - \mathbf{a}\|_p^p\} \text{ and } \nu = \max_{i \in \mathcal{I}^-} \{\|\mathbf{x}_i - \mathbf{a}\|_p^p\}.$$

If we take $\bar{\mathbf{w}}_{1:d+1} = (\mathbf{0}_{1:d}^\top, 1)^\top$, then for any $j \in \mathcal{I}^+$ and $k \in \mathcal{I}^-$ we have

$$\bar{\mathbf{w}}_{1:d+1}^\top \phi_{p,1}(\mathbf{x}_j | \mathbf{a}) = \|\mathbf{x}_j - \mathbf{a}\|_p^p \geq \mu > \nu \geq \bar{\mathbf{w}}_{1:d+1}^\top \phi_{p,1}(\mathbf{x}_k | \mathbf{a}) = \|\mathbf{x}_k - \mathbf{a}\|_p^p.$$

We have just obtained the strict linear separation. The same line of arguments can be followed to show the rest of the proposition after selecting the particular dimension $\ell \in \{1, \dots, d\}$. Suppose that the first part of (4) holds and define

$$\mu_\ell = \min_{i \in \mathcal{I}^+} \{|x_{i\ell} - a_\ell|^p\} \text{ and } \nu_\ell = \max_{i \in \mathcal{I}^-} \{|x_{i\ell} - a_\ell|^p\}.$$

Take $\bar{\mathbf{w}}_{1:2d} = (\mathbf{0}_{1:d}^\top, \mathbf{e}_\ell^\top)^\top$, where \mathbf{e}_ℓ shows the ℓ th unit vector. Then for any $j \in \mathcal{I}^+$ and $k \in \mathcal{I}^-$ we have

$$\bar{\mathbf{w}}_{1:2d}^\top \phi_{p,d}(\mathbf{x}_j | \mathbf{a}) = |x_{j\ell} - a_\ell|^p \geq \mu_\ell > \nu_\ell \geq \bar{\mathbf{w}}_{1:2d}^\top \phi_{p,d}(\mathbf{x}_k | \mathbf{a}) = |x_{k\ell} - a_\ell|^p.$$

We have just obtained the strict linear separation. The second parts of both (3) and (4) can simply be obtained by reversing the roles of \mathcal{I}^+ and \mathcal{I}^- .

S.2 Intermediate dimensions and multi-class case

Proposition 1 provides the conditions for achieving a linear separation by using the proposed feature maps. However, satisfying these conditions with a single anchor point may be difficult. Consider instead the case when a set of anchor points is provided. The following result shows the conditions under which we can obtain a linearly separable dataset. Here, each sample is associated with one of the anchor points in the set. The proof of this result can be shown with the same arguments as in the proof of Proposition 1.

Corollary 1 *In addition to the setup as in Proposition 1, suppose that we also have a set of anchor points \mathcal{A} . If we further define for all $i \in \mathcal{I}$, the feature maps $\phi_{p,1}(\mathbf{x}_i | \mathbf{a}_i)$ with $\mathbf{a}_i = \arg \min_{\mathbf{a} \in \mathcal{A}} \{\|\mathbf{x}_i - \mathbf{a}\|_p\}$, then we obtain a linearly separable dataset when the chosen anchor points satisfy*

$$\min_{i \in \mathcal{I}^+} \{\|\mathbf{x}_i - \mathbf{a}_i\|_p^p\} > \max_{i \in \mathcal{I}^-} \{\|\mathbf{x}_i - \mathbf{a}_i\|_p^p\} \text{ or } \max_{i \in \mathcal{I}^+} \{\|\mathbf{x}_i - \mathbf{a}_i\|_p^p\} < \min_{i \in \mathcal{I}^-} \{\|\mathbf{x}_i - \mathbf{a}_i\|_p^p\}. \quad (5)$$

Using a set of anchor points brings flexibility. For instance, one may simply select \mathcal{A} as all the samples from one of the classes. This selection immediately satisfies the conditions stated in Corollary 1. However, it is important to keep in mind that such a set is likely to lead to overfitting. Therefore, we next introduce another approach, where two sets of anchor points \mathcal{A}_1 and \mathcal{A}_2 are formed with selection of samples from both classes. Then, the feature map that increases the dimension by two for $i \in \mathcal{I}$ simply becomes

$$\phi_{p,2}(\mathbf{x}_i | \mathbf{a}_i^{(1)}, \mathbf{a}_i^{(2)}) = (\mathbf{x}_i, \|\mathbf{x}_i - \mathbf{a}_i^{(1)}\|_p^p, \|\mathbf{x}_i - \mathbf{a}_i^{(2)}\|_p^p)^\top,$$




Figure 4: The discriminant functions obtained with selecting two sets of anchor points from different classes.

where $\mathbf{a}_i^{(1)} = \arg \min_{\mathbf{a} \in \mathcal{A}_1} \{\|\mathbf{x}_i - \mathbf{a}\|_p\}$ and $\mathbf{a}_i^{(2)} = \arg \min_{\mathbf{a} \in \mathcal{A}_2} \{\|\mathbf{x}_i - \mathbf{a}\|_p\}$. Figure 4 gives two illustrations of discriminant functions that are obtained with this feature map.

Then comes the important issue of forming the anchor sets. In our numerical experiments, we have observed that selecting the closest or the farthest sample as the anchor point of a particular sample does not perform well as they are likely to be the support vectors or the outliers (a possible increase in model variance). Thus, one can select the samples within a certain distance that is bounded from above and below. Figure 5 illustrates the distributions of distances that can be used with the new feature map $\phi_{p,2}$. Among these values, we have used only a lower bound to remove the outliers shown in the south-west corners of both plots. Then, the remaining samples are used to form the anchor point sets. We have applied a simple grid search to decide the lower bound. As the reported metrics in each plot show, we have improved our previous accuracies significantly with this simple approach.




Figure 5: The last two components of $\phi_{1,2}$ and the performance metrics. In both tests, we have used the same setup as in our numerical experiments section. The best classification accuracy results that we have obtained before on Fourclass and Phoneme datasets were 79.70% and 78.36%, respectively.

Note that while searching for the anchor points, the training time remains the same with each trial. This is, however, not the case with other maps in the literature, since they all increase the dimension of the feature space. Moreover, the search for anchor points could also be invaluable for the practitioners, since they can gain important insights about the *critical* samples and the structure of their datasets.

Finally, we give a discussion on multi-class classification. Consider a dataset with M classes, *i.e.*, $y_i \in \{1, 2, \dots, M\}$, $i \in \mathcal{I}$. We next introduce for $i \in \mathcal{I}$ the feature map

$$\phi_{p,M}(\mathbf{x}_i | \mathbf{a}_i^{(1)}, \dots, \mathbf{a}_i^{(M)}) = (\mathbf{x}_i, \|\mathbf{x}_i - \mathbf{a}_i^{(1)}\|_p^p, \dots, \|\mathbf{x}_i - \mathbf{a}_i^{(M)}\|_p^p)^\top,$$

where $\mathcal{A}_1, \dots, \mathcal{A}_M$ are the anchor sets and $\mathbf{a}_i^{(j)} = \arg \min_{\mathbf{a} \in \mathcal{A}_j} \{\|\mathbf{x}_i - \mathbf{a}\|_p\}$ for $j = 1, \dots, M$. When there are two classes, we clearly obtain $\phi_{p,2}$. Figure 6 shows a comparison of $\phi_{p,M}$ against two well-known kernel approximation methods, Fourier transform [16] and Nyström [19]. The sets of anchor points are selected with a similar framework that described for the binary problems. As we have observed in Section 3, the proposed multi-class feature map performs better than the approximation methods and gives a classification accuracy on par with RBF.

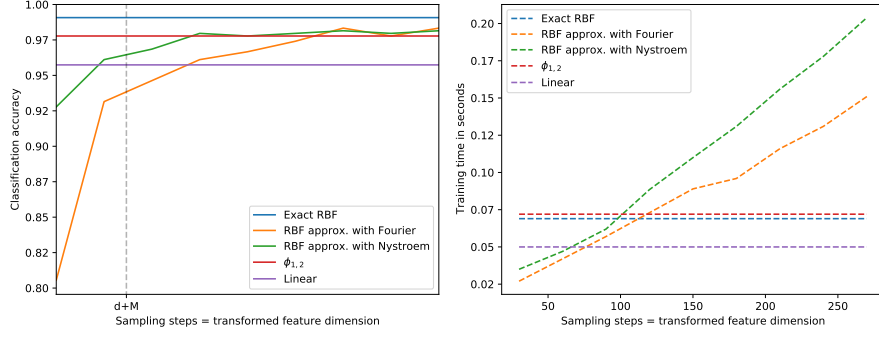


Figure 6: Accuracy and training time comparison against approximation methods on Pen Digits dataset ($M = 10$). This dataset comes with scikit-learn [14].

S.3 Relation to polyhedral conic functions

Our main discriminant functions have ties with so-called conic functions in the literature. The very first work on polyhedral conic functions (PCFs) is given by Gasimov and Ozturk [8]. This work is followed by a series of studies proposing various PCF-based classifiers [13, 3, 6, 5]. In all these works, PCFs are considered as constraints within optimization problems solved for classification. In fact, our current work establishes that the classifiers based on PCFs can be simply considered as kernel methods using explicit feature maps.

Gasimov and Ozturk [8] have proposed the function $g : \mathbb{R}^d \rightarrow \mathbb{R}$ given by

$$g(\mathbf{x}) = \mathbf{w}^\top (\mathbf{x} - \mathbf{a}) + w_{d+1} \|\mathbf{x} - \mathbf{a}\|_1 - b, \quad (6)$$

where $\mathbf{w}, \mathbf{a} \in \mathbb{R}^d$ and $w_{d+1}, b \in \mathbb{R}$. Note that except the displacement with \mathbf{a} in the first term, this function is equivalent to one of our discriminant functions, $f_{1,1}$. The function in (6) is polyhedral conic, since its graph is a cone with vertex at $(\mathbf{a}, -b) \in \mathbb{R}^d \times \mathbb{R}$ and all its sublevel sets are polyhedrons [8, Lemma 2.1]. The authors propose an iterative classification algorithm, which is based on solving a series of linear programming problems. At each iteration, a new vector \mathbf{a} is selected randomly, and the corresponding constraint using function (6) is used to form a new linear programming model. Due to this iterative structure and random selection, the training time of the proposed classifier is not comparable with the state-of-the-art classifiers, and its prediction performance depends on the choice of \mathbf{a} vectors. Later, the authors have also tried clustering methods to select these vectors [13] and considered different ℓ_p -norms [6]. It is important to note that the term *polyhedral conic separation* defined in [8] corresponds to a linear separation in the $d + 1$ dimensional feature space. This is actually the point of view that we advocate in Section 2.

As an extension of [8], Cevikalp and Triggs [3] consider function $h : \mathbb{R}^d \rightarrow \mathbb{R}$ given by

$$h(\mathbf{x}) = \mathbf{w}^\top (\mathbf{x} - \mathbf{a}) + \sum_{\ell=1}^d w_{d+\ell} |x_\ell - a_\ell| - b, \quad (7)$$

where $\mathbf{w}, \mathbf{a} \in \mathbb{R}^d$ and $b, w_{d+\ell} \in \mathbb{R}$ for $\ell = 1, \dots, d$. This time, our discriminant function, $f_{1,d}$ is equivalent to (7) except the displacement with \mathbf{a} in the first term. Unlike others, Cevikalp and Triggs [3] use both (6) and (7) in SVM optimization model and report results on a set of visual object detection problems. In a follow-up study [2], the authors have also taken the squares of the absolute value terms in (7) and obtained ellipsoidal conic functions. They have also considered (6) with ℓ_2 -norm instead of ℓ_1 -norm. In the same study, the authors have also hinted that the samples are “explicitly mapped to a higher feature space.” However, they have not followed with this line of thought to discuss a general framework for the feature maps and the associated kernels as we do here.

As a last note, we point out that none of the work above has considered the intermediate dimensions, nor given a formal discussion on conditions for linear separation (see Proposition 1 and Corollary 1), which allows us to consider a systematic way to explore multiple anchor points.

S.4 Details of the datasets

In this section, we give list the properties of the datasets that we have used in our numerical experiments. All these datasets are also provided with the accompanying compressed file.

Table S.1: Properties of the datasets

Datasets	m	d
Australian	690	14
Fourclass	862	2
Ionosphere	351	34
Heart	270	13
Pima Indians Diabetes	768	8
Wisconsin Prognosis	194	33
Bupa Liver	341	6
Fertility	100	9
Wisconsin Diagnosis	683	10
Splice	3175	60
Wilt	4889	5
SVM Guide1	7089	4
Spambase	4601	57
Phoneme	5404	5
Magic	19020	10
Adult	48852	123